

The AI Readiness Scorecard

Twelve questions across four dimensions. Score each 0, 1, or 2, for a project you can name.

1. DATA FOUNDATION

What the model learned from, and whether anyone checked it.

Do you know the class balance of your training data?

0 Not checked · 1 Known but unadjusted · 2 Resampled or reweighted, with rare-class performance tracked

Has the data been checked for outcome differences across groups?

0 Not yet · 1 Once, at build time · 2 Recurring, reviewed outside the build team

Do you know how your labels were produced?

0 Came with the dataset · 1 Internal, no guidelines · 2 Documented guidelines with measured labeler agreement

2. EXPLAINABILITY

Whether you can account for one decision, not the model in general.

Can you explain a single decision to the person it affected?

0 No · 1 Only in general terms · 2 Per-decision attribution, in language a non-specialist follows

Who has asked you for that explanation so far?

0 Nobody, so it was not built · 1 Internal teams, ad hoc · 2 Built in advance for a regulator, auditor, or customer

Do you know what accuracy a fully interpretable model would cost you?

0 Never compared · 1 Compared once at the start · 2 Re-checked at every retraining

3. MEASUREMENT

Whether you would be able to prove this worked.

What did performance look like before the model existed?

0 Not measured · 1 Rough, informal numbers · 2 Documented baseline on the same metric you now report

Which single metric decides whether this model is working?

0 Accuracy · 1 Accuracy plus one or two others · 2 Chosen for the cost of the error that hurts most

Do you collect signal from the people affected by the output?

0 No · 1 Complaints reach a support inbox · 2 Routed back into evaluation data, reviewed on a schedule

4. GOVERNANCE & MONITORING

What happens on the day the model is wrong.

How would you find out the model has drifted?

0 Someone would notice eventually · 1 Periodic manual review · 2 Automated monitoring, alert thresholds, retraining trigger

Who is accountable when the model is wrong?

0 Unclear · 1 The team that built it · 2 A named owner outside the build team who can pause the system

What personal data does it use, and does it need all of it?

0 Not reviewed · 1 Reviewed at build time · 2 Deliberately minimized, anonymized or federated where practical

READING THE SCORE

0 to 9 Not ready · **10 to 15** Early · **16 to 20** Approaching · **21 to 24** Production-ready

The weakest-link rule. Any dimension scoring 1 or below caps the verdict at Not ready, whatever the total. Any dimension at 3 or below caps it at Approaching. A project is bounded by its weakest dimension, not by its average, because that is how these systems fail: strong modeling does not compensate for unchecked data, and a high total with one dimension near zero is a project with a specific, predictable problem.

Readiness is not permission. A high score says a project is built responsibly. It says nothing about whether it should be built.